

Foundations of Machine Learning (ECE 5984)

- Dimensionality Reduction -

Eunbyung Park

Assistant Professor

School of Electronic and Electrical Engineering

[Eunbyung Park \(silverbottlep.github.io\)](https://silverbottlep.github.io)

Eigenvalues and Eigenvectors

Matrix Decomposition

- We can decompose an integer into its prime factors
 - $12 = 2 \times 2 \times 3$.
- Similarly, matrices can be decomposed into products of other matrices
- Eigendecomposition, SVD, LU decomposition, ...

Eigenvector

- An eigenvector of a square matrix $A \in \mathbb{R}^{n \times n}$ is a nonzero vector v such that

$$Av = \lambda v$$

, where the scalar λ is the eigenvalue

- If v is an eigenvector of A with an eigenvalue λ , then any rescaled αv is also an eigenvectors
- So, usually, we find the '*normalized eigenvectors*'

Compute Eigenvalues

$$Av = \lambda v$$

$$Av - \lambda v = 0$$

$$(A - \lambda I)v = 0$$

- If nonzero solution for v exists, then $(A - \lambda I)$ should be “non-invertible”.

$$\det(A - \lambda I) = 0$$

- A.k.a, characteristic polynomial

Exercise

- What are the eigenvalues and eigenvectors of A ?

$$A = \begin{bmatrix} -3 & 5 \\ 4 & -2 \end{bmatrix}$$

How Many Distinct Eigenvalues?

- An eigenvector of a square matrix $A \in \mathbb{R}^{n \times n}$ is a nonzero vector v such that

$$Av = \lambda v$$

, where the scalar λ is the eigenvalue

How Many Distinct Eigenvalues?


- An eigenvector of a square matrix $A \in \mathbb{R}^{n \times n}$ is a nonzero vector v such that

$$Av = \lambda v$$

, where the scalar λ is the eigenvalue

- There can be maximum distinct 'n' eigenvalues.
- The eigenvalues of an n by n matrix are the roots of a polynomial of degree n. So there are n eigenvalues, though some of them may be repeated.

$$AQ = Q\Lambda \quad A = Q\Lambda Q^{-1}$$



Symmetric Eigendecomposition

- If A is a symmetric (also square) matrix, then
- All the eigenvalues are real
- The eigenvectors corresponding to different eigenvalues are orthogonal
- If we normalize all eigenvectors, then

$$Q Q^T = ?$$

$$A Q = Q \Lambda \quad A = Q \Lambda Q^{-1}$$

Symmetric Eigendecomposition


- If A is a symmetric (also square) matrix, then
- All the eigenvalues are real
- The eigenvectors corresponding to different eigenvalues are orthogonal
- If we normalize all eigenvectors, then

$$Q Q^T = I$$

$$A Q = Q \Lambda \quad A = Q \Lambda Q^{-1} \quad A = Q \Lambda Q^T$$

Symmetric Decomposition


$$A = Q\Lambda Q^{-1} = Q\Lambda Q^{\top}$$



Geometric Interpretation of Eigendecomposition

- Matrix is all about linear transformation!
- Orthogonal matrices \approx Rotational matrices
- $Ax \rightarrow$ scale and rotate the vector x


$$Ax = Q\Lambda Q^T x$$



Principle Component Analysis (PCA)


Dimensionality Reduction

- Redundant features
 - E.g., mph, kph
- Correlation between features
 - E.g., enjoying study, grade, skill





Principal Component Analysis (PCA)

- Finding 'principal' component that explains the data




Maximizing The Variance

- Finding unit vector u , after data projection, the variance of the projected data is maximized



Maximizing The Variance


- Finding unit vector u , after data projection, the variance of the projected data is maximized



$x^{(i)\top} u$: The length of the projection of $x^{(i)}$ onto u

Maximizing The Variance

- Finding unit vector u , after data projection, the variance of the projected data is maximized



$$\begin{aligned}\frac{1}{m} \sum_{i=1}^m \left(x^{(i)\top} u \right)^2 &= \frac{1}{m} \sum_{i=1}^m u^\top x^{(i)} x^{(i)\top} u \\ &= u^\top \left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)\top} \right) u \quad \text{s.t. } \|u\|_2 = 1\end{aligned}$$

Assuming the data is normalized, zero mean

Optimization

- How to optimize it?

$$\max_u u^\top \Sigma u \quad s.t. \|u\|_2 = 1$$

Constraint -> unconstraint

Lagrangian, take the derivative, set it to zero!

Optimization

- How to optimize it?

$$\max_u u^\top \Sigma u \quad s.t. \|u\|_2 = 1$$

Constraint -> unconstraint

Lagrangian, take the derivative, set it to zero!

$$L(u, \lambda) = u^\top \Sigma u - \lambda(u^\top u - 1)$$

Optimization

- How to optimize it?

$$L(u, \lambda) = u^\top \Sigma u - \lambda(u^\top u - 1)$$

$$\frac{\partial L}{\partial u} = \Sigma u - \lambda u = 0$$

$$\Sigma u = \lambda u$$


PCA and Eigenvector

Covariance matrix

Eigenvalue

$$\Sigma u = \lambda u$$

Eigenvector



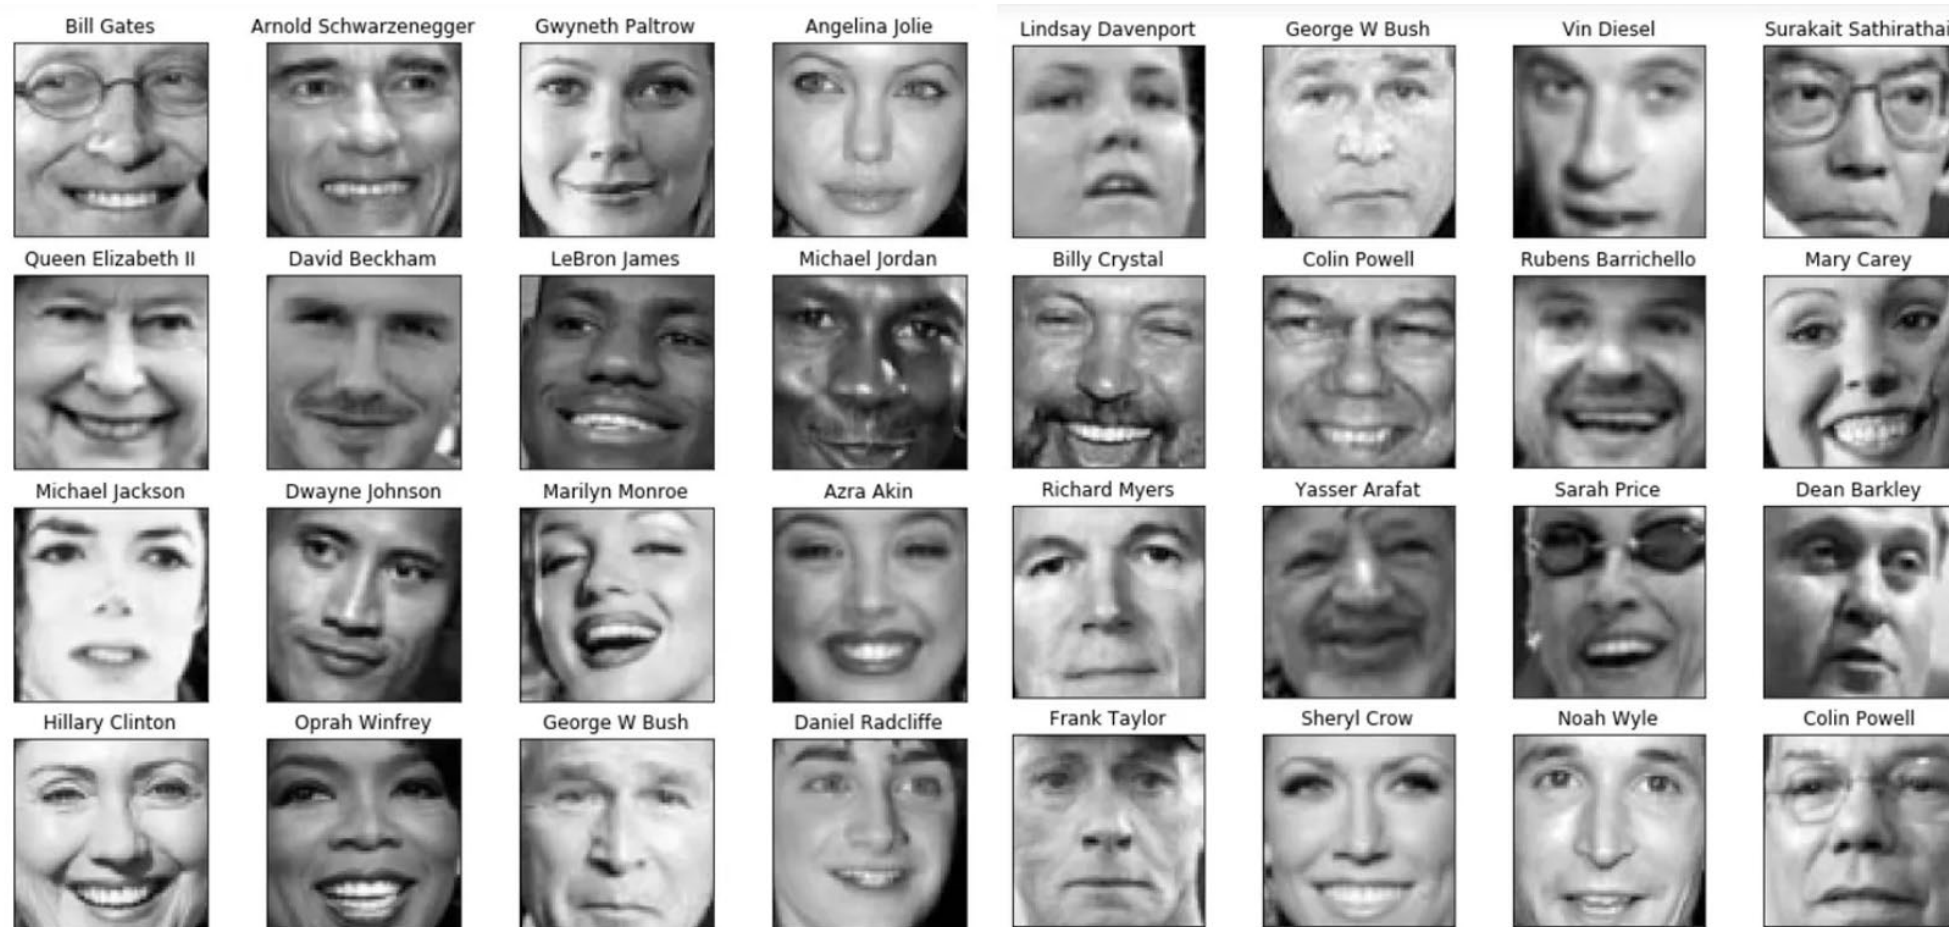
Projected Coordinates

$$y^{(i)} = \begin{bmatrix} u_1^\top x^{(i)} \\ u_2^\top x^{(i)} \\ \vdots \\ u_k^\top x^{(i)} \end{bmatrix} \in \mathbb{R}^k$$

The new coordinate, using top-k
principal component

Eigenface

- Dataset (1000 x 64 x 64 -> 1000 x 4096)




Eigenface

- Dataset (1000 x 64 x 64 -> 1000 x 4096)

$$\Sigma u = \lambda u$$

Eigenface

- Eigenvectors




Code Demo




[05.09-Principal-Component-Analysis.ipynb - Colaboratory \(google.com\)](#)

Nonlinear Dimensionality Reduction

Orthogonal Assumption of PCA




Nonlinear Dimensionality Reduction




Nonlinear Dimensionality Reduction

- Neural Networks based auto-encoder




PCA 2D Embeddings for MNIST

3 6 8 1 7 9 6 6 4 1
6 7 5 7 8 6 3 4 8 5
2 1 7 9 7 1 2 8 4 5
4 8 1 9 0 1 8 8 9 4
7 6 1 8 6 4 1 5 6 0
7 5 9 2 6 5 8 1 9 7
1 2 2 2 2 3 4 4 8 0
0 2 3 8 0 7 3 8 5 7
0 1 4 6 4 6 0 2 4 3
7 1 2 8 7 6 9 8 6 1



t-SNE 2D Embeddings for MNIST



Stochastic Neighbor Embedding (SNE)

- High dimensional neighborhood information as a distribution
- Given $x^{(i)}$, $P_{j|i}$ is the probability that point $x^{(i)}$ chooses $x^{(j)}$ as its neighbor
- Final distribution over pairs is symmetrized

$$P_{j|i} = \frac{\exp\left(-\frac{\|x^{(i)} - x^{(j)}\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{\|x^{(i)} - x^{(k)}\|^2}{2\sigma_i^2}\right)}$$

$$P_{ij} = \frac{1}{2N} (P_{i|j} + P_{j|i})$$

Stochastic Neighbor Embedding (SNE)

- High dimensional neighborhood information as a distribution
- Given $x^{(i)}$, $P_{j|i}$ is the probability that point $x^{(i)}$ chooses $x^{(j)}$ as its neighbor
- Final distribution over pairs is symmetrized

$$P_{j|i} = \frac{\exp\left(-\frac{\|x^{(i)} - x^{(j)}\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{\|x^{(i)} - x^{(k)}\|^2}{2\sigma_i^2}\right)}$$

$$P_{ij} = \frac{1}{2N} (P_{i|j} + P_{j|i})$$

SNE Objective

- Given data, $x^{(1)}, \dots, x^{(N)} \in \mathbb{R}^D$, we define the distribution P_{ij}
- Goal: Find $y^{(1)}, \dots, y^{(N)} \in \mathbb{R}^d$, for some $d \ll D$, minimizing

$$KL(P||Q) = \sum_{ij} P_{ij} \log \left(\frac{P_{ij}}{Q_{ij}} \right) \quad Q_{ij} = \frac{\exp \left(-\|y^{(i)} - y^{(j)}\|^2 \right)}{\sum_{l \neq k} \exp \left(-\|y^{(l)} - y^{(k)}\|^2 \right)}$$

$$P_{j|i} = \frac{\exp \left(-\|x^{(i)} - x^{(j)}\|^2 / 2\sigma_i^2 \right)}{\sum_{k \neq i} \exp \left(-\|x^{(i)} - x^{(k)}\|^2 / 2\sigma_i^2 \right)} \quad P_{ij} = \frac{1}{2N} (P_{i|j} + P_{j|i})$$

KL Divergence

- Measures distance between two distributions, P and Q
- Not a metric function – not symmetric
- $KL(P||Q) \geq 0$
- $KL(P||Q) = 0$ only when $P == Q$

$$KL(P||Q) = \sum_{ij} P_{ij} \log \left(\frac{P_{ij}}{Q_{ij}} \right)$$

Optimizing SNE


$$\begin{aligned}\min_{y^{(1)} \dots y^{(N)}} KL(P||Q) &= \min_{y^{(1)} \dots y^{(N)}} \sum_{ij} P_{ij} \log \left(\frac{P_{ij}}{Q_{ij}} \right) \\ &= \min_{y^{(1)} \dots y^{(N)}} - \sum_{ij} P_{ij} \log(Q_{ij}) + \text{const}\end{aligned}$$

$$\frac{\partial}{\partial y^{(i)}} - \sum_{ij} P_{ij} \log(Q_{ij}) = \dots = \sum_j (P_{ij} - Q_{ij})(y^{(i)} - y^{(j)})$$

- Gradient descent!
- Non-convex, multiple runs!
- Main issue – crowding problem


Crowding Problem

- In high dimension, we have more room
- In low dimension, we do not have enough room to accommodate all neighbors



t-SNE

- t-Distributed Stochastic Neighbor Embedding
- Student's t distribution
- Probability goes to zero much slower than a Gaussian



t-SNE

- t-Distributed Stochastic Neighbor Embedding
- We can now redefine Q_{ij} as
- P_{ij} is same as before

$$Q_{ij} = \frac{\left(1 + \|y^{(i)} - y^{(j)}\|^2\right)^{-1}}{\sum_{l \neq k} \left(1 + \|y^{(l)} - y^{(k)}\|^2\right)^{-1}}$$

t-SNE Algorithms

Algorithm 1: Simple version of t-Distributed Stochastic Neighbor Embedding.

Data: data set $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$,

cost function parameters: perplexity $Perp$,

optimization parameters: number of iterations T , learning rate η , momentum $\alpha(t)$.

Result: low-dimensional data representation $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$.

begin

 compute pairwise affinities $p_{j|i}$ with perplexity $Perp$ (using Equation 1)

 set $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

 sample initial solution $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$ from $\mathcal{N}(0, 10^{-4}I)$

for $t=1$ **to** T **do**

 compute low-dimensional affinities q_{ij} (using Equation 4)

 compute gradient $\frac{\delta C}{\delta \mathcal{Y}}$ (using Equation 5)

 set $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

end

end

t-SNE Visualization

- [Visualizing MNIST: An Exploration of Dimensionality Reduction - colah's blog](#)